

# Wrangling Pre-harmonized Data: Best Practices Based on Harmonizing Income Inequality and Aggregate Public Opinion Measures\*

**Preliminary version. No cycling without permission.**

Frederick Solt<sup>†,1</sup>      Yue Hu<sup>†,2,\*</sup>

This article focuses on a preliminary step in any ex-post data harmonization project—wrangling the pre-harmonized data—and suggests best practices for helping scholars avoid errors in this often-tedious work. To provide illustrations of these best practices, the article uses the examples of harmonizing procedures used to produce the Standardized World Income Inequality Database (SWIID), a widely used database that uses Gini indices from multiple institutions (e.g., UN, Eurostat, OECD, Luxembourg Income Study) to create comparable estimates; and the Dynamic Comparative Public Opinion (DCPO) project, which creates a workflow for harmonizing aggregate public opinion data.

<sup>†</sup> These authors contributed equally to this work.

<sup>1</sup> Department of Political Science, University of Iowa, Iowa City, USA

<sup>2</sup> Department of Political Science, Tsinghua University, Beijing, China

\* Correspondence: [Yue Hu <yuehu@tsinghua.edu.cn>](mailto:yuehu@tsinghua.edu.cn)

A growing quantity of political science research can be characterized as data science: it employs large quantities of data, often drawn from a large number of different sources. For such projects, data wrangling, the task of getting these data into the format required to perform analyses, is notoriously the bulk of the work (see, e.g., [Lohr 2014](#)). Such data ‘janitor work’ is often viewed as tiresome, as something to be delegated to research assistants, to someone—indeed anyone—else (see [Torres 2017](#)). Data wrangling is, however,

---

\*Yue Hu appreciates the funding support from the National Natural Science Foundation of China (72374116) and Tsinghua University Initiative Scientific Research Program (2024THZWJC01).

critically important to scientific inquiry, and errors in the process can undermine our conclusions.

We focus here on one particularly insidious problem that can affect the ‘janitor work’ of any researcher: data-entry errors. Faced with the task of getting data into the correct format, even some very sophisticated researchers will conclude that the most straightforward means to that end is to simply copy the needed data into a spreadsheet manually. This technique may be straightforward, but it is very much prone to error. Barchard and Pace (2011a) found that ‘research assistants’ assigned in an experiment to carefully enter data manually, even those instructed to double-check their entries against the original, had error rates approaching 1% in just a single roughly half-hour session. Rates likely go up as the tedious task goes on.

Like ‘janitor work’ itself, data-entry errors have thus far gained little attention in political science. In this piece, we illustrate the pernicious threat this problem poses by carefully scrutinizing a prominent recent work that examines how changes in democracy affect democratic support among the public (Claassen 2020a).<sup>1</sup> We document the data-entry errors that slipped past both the author and the journal’s strict replication policy and how these errors affect the paper’s results and conclusions. Claassen (2020a, 51) concludes that when “elected leaders start dismantling democratic institutions and rights, public mood is likely to swing rapidly toward democracy again, providing something of an obstacle to democratic backsliding.” We show that, after data-entry errors are corrected, there is no empirical evidence that public support responds thermostatically to changes in democracy in this way.

Before elaborating, we note that it is impossible to tell with complete certainty the exact reason for the issues we identify; it is possible that coding mistakes or even intentional decisions are at fault rather than data-entry errors strictly speaking. Further, our point here is not to criticize a particular result, but to highlight a case that “reflects on typical robustness challenges” (Janz and Freese 2021, 306) and so to illuminate how idiosyncratic manual data entry processes can cause problems for empirical research. On this basis,

---

<sup>1</sup>In this study, we leverage the case of thermostatic support for democracy to emphasize the impact of data preprocessing—a largely underexplored but critical phase in research methodology—on empirical outcomes. As Weidmann (2023) points out, although data collection, processing, and analyses can all influence conclusions, they follow distinct principles and require different types of action. Each of them merits specific attention, especially the data processing stage, our focus here (Weidmann 2023, 4–9). On the impact of data analysis, see (?), which provides an in-depth discussion on that issue in the context of a similar substantive topic.

we conclude with four practical suggestions to help political scientists reduce data-entry errors and their impact. Rather than merely a single replication of a difficult-to-detect phenomenon, we hope readers will also consider this work as a contribution to the growing “open science” movement to build more reliable and transparent research both in and beyond social science (Christensen, Freese, and Miguel 2019; Engzell and Rohrer 2021).

## 0.1 Data-Entry Errors and Democratic Support

With democracy under increasing threat in countries around the world, how the public reacts is a crucial question. According to a classic and still vibrant literature, growing experience with democratic governance helps generate robust public support for democracy (Welzel 2013; Wuttke, Gavras, and Schoen 2022; see, e.g., ?). Claassen (2020a) argues instead that democratic support behaves thermostatically: that increases in democracy yield an authoritarian backlash in the public, while democratic backsliding prompts the public to rally to democracy’s cause.

The evidence it offers in support of this latter argument takes advantage of recent advances in modeling public opinion as a latent variable to measure democratic support. This approach provides estimates of the paper’s dependent variable for over one hundred countries for up to nearly three decades, constituting a much larger evidentiary base than any previous study. These latent-variable estimates, in turn, were based on thousands of nationally aggregated responses to dozens of different questions from cross-national survey projects (Claassen 2020a, 40). Two pieces of data were collected for each distinct survey item in each country and year it was asked: the number of respondents to give a democracy-supporting response—defined, for ordinal responses, as those above the median value of the scale (Claassen 2020a, Appendix 1.3)—and the total number of respondents to whom the question was posed. Each of these 7,538 pieces of source data is recorded in a spreadsheet.<sup>2</sup>

We re-collected all of the source data for the publication from the original surveys. We identified the variables of the survey items used by the article within each survey dataset, and then we used an automated process to collect the needed data from the survey

---

<sup>2</sup>The article’s replication materials include only the latent variable estimates without the original survey aggregates that serve as their source data (see Claassen 2020c). Fortunately, however, the spreadsheet recording these original source data is included in the replication materials for a companion piece that employed the identical estimates (see Claassen 2020b).

datasets while avoiding data-entry errors (see [Solt, Hu, and Tai 2018](#)). In Figure 1, we compare the percentage of respondents to give a democracy-supporting response in the publication spreadsheet with the percentage we found using our automated process of wrangling these same data.




Figure 1: Comparing Democracy-Supporting Responses in the Publication Data and the Corrected Data. The x-axis represents the corrected data, and the y-axis represents the publication data. The points are colored by the item family and shaped by the project. The dotted line represents the line of equality. The plot shows a strong positive correlation between the two variables, with most points falling above the line of equality.

For 85% of the country-year-item observations, the difference between these percentages was negligible—less than half a percent—yielding points approximately along the plot’s dotted line. But for the remaining observations, the difference was often substantial due to data-entry errors in the publication data. For example, the Asia Barometer asked respondents in 35 country-years to indicate whether they thought “a democratic political system” would be very good, fairly good, or bad for their country. According to the study’s coding rules (see [Claassen 2020a](#), Appendix 1.3), only answers above the median of the response categories should be considered as democracy supporting, yet in this case the lukewarm intermediate category was coded as supporting democracy as well.<sup>3</sup> This

<sup>3</sup>Although this may be interpreted as an exercise of researcher judgment as to what constitutes a democracy-supporting response rather than a data-entry error, examination of similar answers to similar questions shows that similarly lukewarm responses at and below the median response category

led to overestimations of the percentage of democracy-supporting responses ranging from 19 to 63 percentage points and averaging 44 points.

Similarly, the four waves of the Asian Barometer included the following item: “Here is a similar scale of 1 to 10 measuring the extent to which people think democracy is suitable for our country. If 1 means that democracy is completely unsuitable for [name of country] today and 10 means that it is completely suitable, where would you place our country today?” In accordance with the coding rules of the study, responses of 6 through 10 are considered democracy supporting, and that is how the first, third, and fourth waves of the survey are coded. For the second wave, however, 5 was erroneously also included among the democracy-supporting responses. This data-entry error resulted in overestimates of as much as 23 percentage points in 9 country-years.

A third example comes from the Pew Global Attitudes surveys’ four-point item asking about the importance of living in a country with regular and fair contested elections: the question wording is “How important is it to you to live in a country where honest elections are held regularly with a choice of at least two political parties? Is it very important, somewhat important, not too important or not important at all?” In this case, rather than including respondents who gave both responses above the median—“very important” and “somewhat important”—only those respondents who answered “very important” were entered as supporting democracy. This error caused substantial underreporting of the extent of democratic support in 91 country-years.

While these issues involve mistakes in recording the numerator of the percentage, the number of respondents who provided a democracy-supporting answer, entering the denominator, the total number of respondents asked a question, was also problematic on occasion. For example, when the Americas Barometer surveyed Canada in 2010, asked half its sample, when “democracy doesn’t work,” Canadians “need a strong leader who doesn’t have to be elected through voting.” Those who were not asked the question were included in the total number of respondents. According to the study’s coding rules, refusing to answer is equivalent to answering in a fashion not supporting democracy (see [Claassen 2020a](#), Appendix 1.3). This rule may or may not be a reasonable coding choice, but including in this category those who were never asked the question at all is clearly a data-entry error.

---

(e.g., in the Arab Barometer, that democracy was “somewhat appropriate” for the country) were coded as not supportive.

Another source of data-entry errors here involves survey weights. Weighting raw survey results to maximize the extent to which they are representative of the target population is important. Relying on topline reported in codebooks rather than the survey data itself evidently caused some mistakes in correctly entering the needed information here, as codebooks do not always take survey weights into account. These errors shifted the percentage of democracy-supporting responses in both directions, typically by relatively small amounts.

Finally, although not depicted on this plot, data-entry errors were also evident in the variable recording the year in which a survey was conducted: these typically reflected differences between the nominal year of a survey wave and when the survey was actually in the field in a particular country. This was an issue for some 9% of the country-year observations.

## 0.2 Discussion

The analysis above reveals that data-entry errors are an especially pernicious threat to the credibility of our results. The threat is a subtle one that is not easily detected. To discern it requires close scrutiny of every manual entry; merely examining the data and their distribution will uncover few errors ([Barchard and Pace 2011a, 1837–38](#)). Although failure to find support for a research hypothesis may prompt us to undertake a such a close review, an analysis that yields statistical significance is unlikely to trigger what will likely be, as in the above example, a time-consuming and difficult effort (see [Gelman and Loken 2014, 464](#)).

This leads us to recommend the following steps to reduce possible data-entry errors. First, **automate data entry**: we suggest researchers to consider reducing reliance on manual data entry and increase the extent to which data wrangling is performed computationally. Readers should be aware that our suggestion does *not* imply that computers are always superior to people for this purpose. Instead, automated coding should always involve sensitive human design and systematic supervision (see [Brezna 2021](#); [Grimmer 2015](#)). Still, given the convenience of automatic data entry in documentation and reproduction, we encourage researchers to use it as much as possible instead of entering data manually to increase the efficacy and transparency of their data processing operations.<sup>4</sup>

---

<sup>4</sup>For a systematic discussion of the function of automated data processes in social science, see Weidmann

In making this recommendation, we are aware that being open and transparent in this way takes effort ([Engzell and Rohrer 2021](#)). But as researchers automate more of their data entry, the chances that they can reuse their code in subsequent projects improve. In fact, many common janitor-work chores already have been packaged as open-source software that to make researchers' task more straightforward and easier.<sup>5</sup>

Second, **use the double-entry method**: when manual data entry cannot be avoided, each entry should be made twice. Double entry is labor intensive, but experiments have shown that it reduces error rates by thirty-fold even when done immediately after the initial collection and by the same person ([Barchard and Pace 2011a, 1837](#)). Given that data-entry errors can completely undermine the validity of our conclusions, as in the example above, double entry is worth the extra effort.

Third, **embrace teamwork**: for any project involving entering data by hand, splitting the task up among team members will reduce the risk of errors going undetected. When double entries are performed by different people, discrepancies will be noted, discussed, and resolved correctly; having two sets of eyes on complex materials like survey codebooks also increases the chances that nuances of the presentation like survey weights will be uncovered. Further, by dividing the load, teamwork also lessens the probability of errors due to fatigue arising in the first place.

Fourth, **be aware of the threat of data-entry error**: this final recommendation is especially for manuscript reviewers. If data-entry errors are invisible to the authors themselves, they are doubly so to reviewers (though if editors provided reviewers with replication materials at the time of the review it may help them to better assess the work's credibility). But the case described above nevertheless suggests a valuable heuristic: when a work's conclusions suggest that a difficult problem will be easily solved—that democratic erosion will reflexively trigger a backlash and a renewed public support for democracy, in the present instance—it warrants especially careful scrutiny.

Data-entry errors are inevitable, and even following these recommendations is unlikely to eliminate them entirely. Further, the above suggestions follow closely from a specific case and, although they successfully help us identify and fix its data-entry issues, they do not constitute a panacea to cure all data-processing problems in all types of research.

---

(2023).

<sup>5</sup>For example, see `readtext` ([Benoit et al. 2016](#)) for formatting text files and `DCP0tools` ([Solt, Hu, and Tai 2018](#)) for aggregating cross-sectional time-series public-opinion surveys.

Nonetheless, we also hope the readers to see the shared logic of these suggestions and the growing literature to guide political scientists to conduct more reliable and credible research. For instance, in the same vein as our first suggestion, Weidmann (2023) provides a book-length set of illustrations on how to reduce “manual point and click” tasks found in a variety of studies with the `tidy-data` framework in the R language. Kapiszewski and Karcher (2021, 288) even suggests that qualitative researchers should consider using “open-exchange format” of qualitative data analysis software to be more “transparent about the generation and analysis of data.” Furthermore, we regard our efforts and recommendations as a contribution to the open science movement to produce more robust and credible research in the social sciences (see, e.g., Christensen, Freese, and Miguel 2019) and beyond (see, e.g., Barchard and Pace 2011b; Lohr 2014). With careful attention, not only can the threat of data-entry errors to our ‘janitor work’, our research, and our understanding of the world be minimized, but the transparency, openness, and credibility of our research can continuously grow.

## Reference

- Barchard, Kimberly A., and Larry A. Pace. 2011. “Preventing Human Error: The Impact of Data Entry Methods on Data Accuracy and Statistical Results.” *Computers in Human Behavior* 27(5): 1834–39.
- Benoit, Kenneth, Adam Obeng, Paul Nulty, Aki Matsuo, Kohei Watanabe, and Stefan Müller. 2016. “readtext: Import and Handling for Plain and Formatted Text Files.”
- Claassen, Christopher. 2020a. “In the Mood for Democracy? Democratic Support as Thermostatic Opinion.” *American Political Science Review* 114(1): 36–53.
- Claassen, Christopher. 2020b. “Replication Data for: Does Public Support Help Democracy Survive?”
- Claassen, Christopher. 2020c. “Replication Data for: In the Mood for Democracy? Democratic Support as Thermostatic Opinion.”



- Gelman, Andrew, and Eric Loken. 2014. “The Statistical Crisis in Science.” *American Scientist* 102(6): 460–65.
- Solt, Frederick, Yue Hu, and Yuehong ‘Cassandra’Tai. 2018. “DCPOtools: Tools for Dynamic Comparative Public Opinion.” <https://github.com/fsolt/DCPOtools>.
- Welzel, Christian. 2013. *Freedom Rising: Human Empowerment and the Quest for Emancipation*. Cambridge: Cambridge University Press.
- Wuttke, Alexander, Konstantin Gavras, and Harald Schoen. 2022. “Have Europeans Grown Tired of Democracy? New Evidence from Eighteen Consolidated Democracies, 1981–2018.” *British Journal of Political Science* 52(1): 416–28.

# Online Supplementary Materials

## A Descriptive Statistics

### Reference

- Barchard, Kimberly A., and Larry A. Pace. 2011. "Preventing Human Error: The Impact of Data Entry Methods on Data Accuracy and Statistical Results." *Computers in Human Behavior* 27(5): 1834–39.
- Breznau, Nate. 2021. "I Saw You in the Crowd: Credibility, Reproducibility, and Meta-Utility." *PS: Political Science & Politics* 54(2): 309–13. doi:[10.1017/S1049096520000980](https://doi.org/10.1017/S1049096520000980).
- Christensen, Garret, Jeremy Freese, and Edward Miguel. 2019. *Transparent and Reproducible Social Science Research: How to Do Open Science*. Berkeley: University of California Press.
- Claassen, Christopher. 2020. "Replication Data for: In the Mood for Democracy? Democratic Support as Thermostatic Opinion."
- Engzell, Per, and Julia M. Rohrer. 2021. "Improving Social Science: Lessons from the Open Science Movement." *PS: Political Science & Politics* 54(2): 297–300. doi:[10.1017/S1049096520000967](https://doi.org/10.1017/S1049096520000967).
- Grimmer, Justin. 2015. "We Are All Social Scientists Now: How Big Data, Machine Learning, and Causal Inference Work Together." *PS: Political Science & Politics* 48(1): 80–83. doi:[10.1017/S1049096514001784](https://doi.org/10.1017/S1049096514001784).
- Janz, Nicole, and Jeremy Freese. 2021. "Replicate Others as You Would Like to Be Replicated Yourself." *PS: Political Science & Politics* 54(2): 305–8. doi:[10.1017/S1049096520000943](https://doi.org/10.1017/S1049096520000943).

Kapiszewski, Diana, and Sebastian Karcher. 2021. “Transparency in Practice in Qualitative Research.” *PS: Political Science & Politics* 54(2): 285–91. doi:[10.1017/S1049096520000955](https://doi.org/10.1017/S1049096520000955).

Lohr, Steve. 2014. “For Big-Data Scientists, ‘Janitor Work’ Is Key Hurdle to Insights.” *The New York Times: Technology*.

Torres, Rachel. 2017. “Me: Shouldn’t There Be Someone in a Basement That We Just Pay to Do All This Awful Data Cleaning? Advisor: That’s Who You Are.”

Weidmann, Nils B. 2023. *Data Management for Social Scientists: From Files to Databases*. Cambridge: Cambridge University Press. doi:[10.1017/9781108990424](https://doi.org/10.1017/9781108990424).